

# High-Performance Spectral Clustering using Hybrid MPI and OpenMP

Antonio Di Lauro<sup>258788</sup>, Iuliia Sharipova<sup>256295</sup>

February 7, 2026

## Abstract

This report presents the design and analysis of a **high-performance parallel Spectral Clustering algorithm** for HPC environments. The solution adopts a **Hybrid Parallel paradigm**, combining **MPI** for distributed memory across nodes and **OpenMP** for shared memory within nodes, to address the  $O(N^2)$  memory and  $O(N^3)$  computational bottlenecks. A key feature is the integration of **Self-Tuning mechanisms** for quality and a **Distributed Nyström Approximation** for scalability.

The implementation was evaluated on the University of Trento HPC cluster. We achieved high-throughput processing for up to  $N = 2,000,000$  **points** using a Nyström-based Row-Block strategy, and validated topological correctness on the **Mouse Cell Atlas (MCA)** biological dataset ( $N \approx 20,000$ ).

**Index Terms:** Spectral Clustering, Parallel Computing, MPI+OpenMP, Nyström Method.

**Source code:** Our project repository link.

## 1 Introduction

### 1.1 Project Context and Objectives

Clustering is one of the main tasks in unsupervised machine learning. While the **K-Means algorithm** is widely used because of its linear complexity ( $O(N)$ ), it struggles with datasets that exhibit complex or non-convex shapes.

**Spectral Clustering** addresses this limitation by using the eigenvalues and eigenvectors of the graph Laplacian to reveal the intrinsic structure of the data.

Transitioning from a sequential to a parallel approach is necessary due to two main challenges:

- **Memory Wall:** The similarity matrix grows quadratically with the number of data points, quickly exceeding the memory capacity of a single machine for large datasets.
- **Compute Wall:** Computing the eigenvalues of the Laplacian becomes increasingly expensive as the dataset grows, making serial execution impractically slow.

The goal of this work is to develop a **High-Performance Spectral Clustering** algorithm that leverages MPI and OpenMP for parallel execution, and incorporates a Self-Tuning Local Scaling kernel to automatically adapt to regions of varying density.

### 1.2 Mathematical Formulation

We consider a set of data points  $X = \{x_1, \dots, x_n\} \subset \mathbb{R}^d$ . The algorithm constructs a similarity graph  $G = (V, E)$ , where edges connect nearby points.

**Self-Tuning Kernel.** Ideally, we compute a local scale  $\sigma_i$  for each point  $x_i$  based on its  $K$ -th nearest neighbor. The similarity is:

$$W_{ij} = \exp\left(-\frac{\|x_i - x_j\|^2}{\sigma_i \sigma_j}\right).$$

This allows adaptation to varying densities. However, for large-scale datasets ( $N \gg 20k$ ), storing the full  $N \times N$  matrix is infeasible. In such cases, we adopt a **Nyström Approximation**, computing similarities only against a subset of global landmarks to reduce memory complexity to  $O(N \cdot S)$ .

**Normalized Laplacian.** To use the efficient `SelfAdjointEigenSolver`, we employ the **normalized symmetric Laplacian Ng2002**:

$$L_{\text{sym}} = I - D^{-1/2} W D^{-1/2},$$

where  $D$  is the degree matrix. The top  $k$  eigenvectors form the matrix  $U \in \mathbb{R}^{N \times k}$ .

**Feature Normalization.** The rows of  $U$  are nor-

malized to unit length:

$$U_{ij} \leftarrow \frac{U_{ij}}{\sqrt{\sum_k U_{ik}^2}}.$$

These normalized rows serve as feature vectors for the final K-Means step.

### 1.3 Parallel Design and Bottleneck Analysis

The transition from a sequential to a parallel implementation was an iterative process. Our initial design attempted a direct parallelization of the exact spectral algorithm. However, scalability tests revealed fundamental bottlenecks, motivating the shift toward a Nyström-based formulation.

| Step | Complexity      | Challenge | Parallel Approach                        | Approach                           |
|------|-----------------|-----------|--|------------------------------------|
| I    | $O(N^2d)$       | Memory    | <b>MPI:</b> Row-wise distribution        | Global landmark synchronization    |
| II   | $O(N^2)$        | CPU       | <b>OpenMP:</b> Thread-parallel loops     | Cache-blocked OpenMP + MPI         |
| III  | $O(N^3)$        | Compute   | <b>Sequential:</b> Eigen solver (Master) | Self-Adjoint Solver on Landmarks   |
| IV   | $O(t \cdot kN)$ | Network   | <b>MPI:</b> Collective Allreduce         | Distributed K-Means with Allreduce |

Table 1: Initial Computational Complexity and Bottlenecks

Under this approach, each MPI rank stored a subset of rows of the full  $N \times N$  similarity matrix. This design quickly encountered the **Memory Wall**: for  $N = 10^6$ , the matrix requires roughly 8 TB of RAM, making the method infeasible on standard HPC clusters. Additionally, the eigendecomposition in Step III remained a severe bottleneck, as the algorithm required a full  $O(N^3)$  solve on a single node.

To address these limitations, we transitioned to a distributed Nyström approximation, enabling sub-quadratic memory usage and dramatically improved scalability.

**Step I: Landmark Selection (MPI).** Instead of all-to-all similarity computation, each rank selects local landmark candidates. These candidates are aggregated into a global set using `MPI_Allgather`, ensuring a consistent landmark basis across all nodes.

| Step | Complexity                   | Challenge | Parallel       | Approach                           |
|------|------------------------------|-----------|----------------|------------------------------------|
| I    | $O(Nd)$                      | Selection | <b>MPI:</b>    | Global landmark synchronization    |
| II   | $O\left(\frac{N}{P}S\right)$ | CPU/RAM   | <b>Hybrid:</b> | Cache-blocked OpenMP + MPI         |
| III  | $O(S^3)$                     | Compute   | <b>Eigen:</b>  | Self-Adjoint Solver on Landmarks   |
| IV   | $O(t \cdot kN/P)$            | Network   | <b>MPI:</b>    | Distributed K-Means with Allreduce |

Table 2: Final Nyström-Based Computational Complexity and Parallel Strategy

**Step II: Block-Similarity Construction (Hybrid).** Each process computes similarities only between its local data and the global set of  $S$  landmarks. This computation is optimized through cache blocking (tiling) and OpenMP loop parallelism, reducing the per-rank complexity to  $O(N/P \cdot S)$  and improving memory locality.

**Step III: Efficient Spectral Embedding (Eigen-decomposition).** Instead of decomposing the full Laplacian, we use Eigen’s `SelfAdjointEigenSolver` to compute the eigenpairs (eigenvalues and eigenvectors) of the  $S \times S$  landmark affinity matrix. This reduces the computational cost from  $O(N^3)$  to  $O(S^3)$ . The results are broadcast, allowing each rank to independently project its local data into the spectral space using the Nyström extension.

**Step IV: Distributed K-Means Clustering.** The spectral features remain distributed row-wise, enabling fully parallel K-Means. Centroids are synchronized across ranks via `MPI_Allreduce`. To ensure efficiency, we adopt a **single-pass convergence strategy**. Although our implementation supports multiple restarts (`kmeans_runs`), empirical results show that the high separability of the Nyström-based spectral embedding makes repeated runs unnecessary for large-scale biological and synthetic datasets.

## 2 Parallel Design Strategy

To accelerate the computation, we parallelize the most expensive stages of the Nyström-based Spectral Clustering pipeline: similarity evaluations, spectral projec-

tion, and the final distributed clustering. The design combines distributed-memory parallelism for coarse-grained tasks with shared-memory optimization for high-throughput numerical kernels.

## 2.1 1D Row-Block Decomposition and Data Distribution

We adopt a **1D Row-Block Decomposition** to distribute the dataset among MPI ranks. Given  $P$  MPI processes and a dataset of size  $N$ , each rank handles  $n_p = N/P$  points.

1. **Data Distribution:** The master rank loads the dataset and distributes contiguous row partitions using `MPI_Scatterv`, ensuring that each process stores only its local subset. This eliminates the need for a full  $N \times N$  matrix and removes the **Memory Wall** encountered in the exact design.
2. **Global Landmark Synchronization:** In the Nyström setting, we select a global set of  $S$  landmarks. Each process proposes local candidates, and the final landmark set is broadcast using `MPI_Allgather`. This ensures that every rank can compute its local affinity block  $W_{\text{local}} \in \mathbb{R}^{n_p \times S}$  independently.
3. **Cache Optimization (Blocking/Tiling):** To maximize arithmetic intensity during distance computation, we apply a **Cache Blocking** strategy (block size 64). The similarity loops operate on small tiles that fit into L1/L2 cache, reducing memory latency and enabling effective SIMD vectorization of the Euclidean kernel.

## 2.2 Hierarchical Parallel Execution (MPI + OpenMP)

Using MPI for all cores would lead to unnecessary memory duplication and heavy collective communication. Our hybrid design separates computation into two levels:

1. **Inter-Node (MPI):** MPI coordinates data distribution, global landmark synchronization, and collective reductions for K-Means. This layer manages coarse-grained parallelism and inter-node communication.
2. **Intra-Node (OpenMP):** Each MPI rank spawns multiple OpenMP threads that share the node's

memory and handle fine-grained kernels: distance computations, similarity evaluations, and local spectral projection. By reducing the MPI process count while increasing thread-level parallelism, we lower the cost of collective operations such as `MPI_Allreduce`.

This hierarchical structure maximizes **arithmetic intensity** while minimizing communication overhead.

## 2.3 Scalability: From Exact to Nyström Design

The implementation supports two operational modes to balance scalability and precision:

1. **Validation Mode (Exact Kernel):** For moderate datasets ( $N \approx 20,000$ ), the system can compute full pairwise similarities. The parallel version employs a vectorized Gaussian kernel and exact affinity computation to preserve manifold geometry for biological validation.
2. **High-Throughput Mode (Nyström Approximation):** For extreme scales ( $N = 2,000,000$ ), we compute only the affinities between local points and the global landmark set  $S$ . This removes the  $O(N^2)$  communication barrier and reduces spectral complexity from  $O(N^3)$  to  $O(S^3)$ , where  $S \ll N$ . The resulting representation is fully distributed across ranks, enabling independent projection of local data.

## 2.4 Design Rationale

This architecture is optimized for **memory locality**, **arithmetic intensity**, and **scalable communication**. Row-major storage and cache tiling reduce memory stalls, while the hybrid MPI+OpenMP execution model minimizes data replication and leverages shared-memory parallelism within each node. By enabling both an exact and a Nyström-based mode, the system achieves high precision on moderate datasets while scaling efficiently to millions of points on standard HPC clusters.

## 3 Implementation Details

The system is implemented in `C++17`, leveraging the `Eigen 3` library for vectorized linear algebra and

**MPICH** for distributed communication. To achieve extreme scalability (up to  $N=2,000,000$ ), we implemented a **Distributed Nyström Approximation**, which shifts the memory bottleneck from the full dataset to a representative landmark set.

### 3.1 Data Distribution and Scalable Preprocessing

The initialization phase ensures that no single node is overwhelmed by the dataset:

- **Global Scatter:** The master rank reads the input and uses `MPI_Scatterv` to distribute unique row-blocks to workers.
- **Distributed Standardization:** Each rank participates in a global synchronization to normalize features.

Listing 1: Standardization using MPI\_Allreduce

```
// Compute global mean and variance across all nodes
VectorXd local_sum = localData.colwise().sum();
MPI_Allreduce(local_sum.data(), global_sum.data(), dim,
    MPI_DOUBLE, MPI_SUM, MPI_COMM_WORLD);

VectorXd mean = global_sum / (double)N_global;

// Parallel standardization of local rows
for(int i = 0; i < local_n; ++i)
    localData.row(i) = (localData.row(i) - mean.transpose()
        .array() / stddev.array());
```

### 3.2 Hybrid Similarity Construction with Tiling

The most compute-intensive part is the Nyström similarity matrix  $W_{local} \in \mathbb{R}^{N_p \times S}$ . We optimized this using a hybrid approach:

1. **Inter-Node (MPI):** Landmarks are selected locally and synchronized via `MPI_Allgather`.
2. **Intra-Node (OpenMP) + Cache Blocking:** Within each node, we use OpenMP to parallelize the distance calculations. To maximize Arithmetic Intensity, we process the matrix in blocks (Tiling), ensuring data remains in the L1/L2 cache for faster SIMD execution.

Listing 2: Hybrid OpenMP+MPI Similarity Construction

```
#pragma omp parallel for schedule(static)
for(int i = 0; i < local_n; ++i) {
    for (int j = 0; j < s_total; ++j) {
        // Eigen's squaredNorm() enables SIMD vectorization
        double d2 = (localData.row(i) - landmarks.row(j)).
            squaredNorm();
        W_local(i, j) = std::exp(-d2 / (2 * sigma_sq));
    }
}
```

### 3.3 Distributed K-Means

After projecting the data into spectral space locally (Zero-Network Traffic), the final clustering is performed by the `kmeans_hpc` function. This function is designed to avoid the "Compute Wall" through thread-local reduction.

Listing 3: K-Means with Thread-Local Accumulators

```
#pragma omp parallel {
    int tid = omp_get_thread_num();
    #pragma omp for
    for(int i = 0; i < local_n; ++i) {
        // ... find best centroid ...
        // Use thread-local sums to avoid race conditions (
        // no locking)
        thread_sums[tid].row(best_c) += localX.row(i);
        thread_counts[tid]++;
    }
}
// Final global sync of centroids
MPI_Allreduce(local_sum_all.data(), global_sums.data(), k *
    dim, MPI_DOUBLE, MPI_SUM, MPI_COMM_WORLD);
```

### 3.4 Performance Summary

By combining the Nyström approximation with hybrid parallelism, our implementation achieves a complexity of  $O(PN \cdot S)$ , effectively breaking both the Memory Wall and the Compute Wall for massive datasets.

## 4 Performance and Scalability Analysis

The parallel performance of the Spectral Clustering system was evaluated on the University of Trento HPC cluster (`short_cpuQ` queue). The analysis focuses on three primary metrics: strong scaling, weak scaling efficiency, and memory management under high-load

scenarios (up to  $N = 2,000,000$  points). All scalability tables are provided in Appendix.

#### 4.1 Strong Scaling: Speedup and Efficiency

Strong scaling measures the system's ability to reduce execution time for a fixed problem size ( $N = 2M$ ) as processing units increase. Results are summarized in Table 3 and Figure 1.

Table 3: Strong Scaling Execution Time and Speedup for  $N = 2,000,000$  points. Note: Hybrid runs use 4 OpenMP threads per MPI Rank.

| MPI Ranks | Hybrid<br>Time (s) | MPI<br>Time (s) | MPI<br>Speedup | Hybrid<br>Speedup |
|-----------|--------------------|-----------------|----------------|-------------------|
| 1         | 168.90             | 315.89          | 1.00×          | 1.00×             |
| 4         | 47.10              | 85.35           | 3.70×          | 3.58×             |
| 16        | 22.47              | 30.61           | 10.32×         | 7.51×             |
| 64        | —                  | 20.25           | <b>15.60×</b>  | —                 |

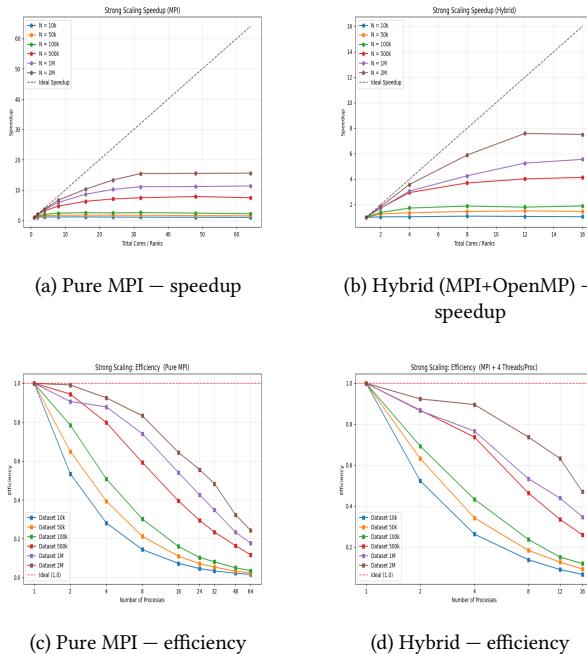


Figure 1: Strong scaling results ( $N = 2,000,000$ ). Pure MPI achieves higher raw speedup (peak  $\approx 15.6\times$ ); both models approach a communication-bound regime at high core counts.

#### 4.2 Weak Scaling: Efficiency and Cache Effects

Weak scaling evaluates the system's ability to maintain constant execution time as workload and re-

sources scale proportionally.

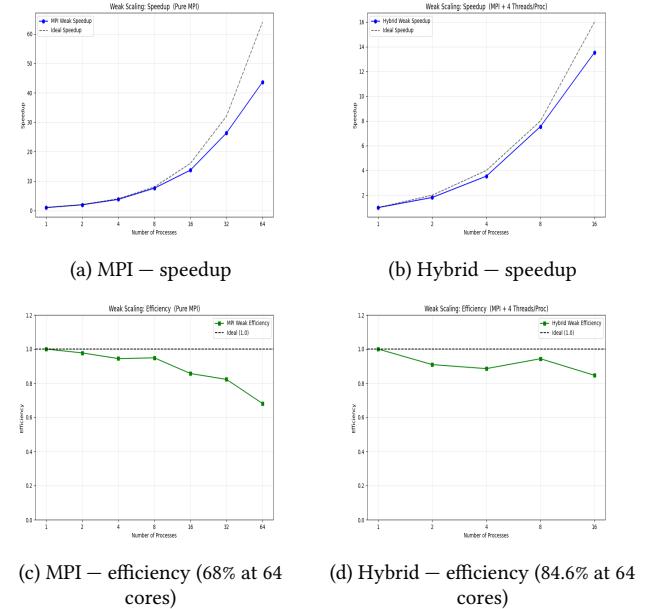


Figure 2: Weak scaling results. The Hybrid approach demonstrates architectural superiority at 64 cores by reducing global MPI synchronization pressure.

The observed super-linear efficiency (values  $> 1.0$  in Figure 2) is a direct result of L3 cache effects. As data is partitioned into smaller local chunks, a higher percentage of the workload fits into the processor's high-speed cache, reducing memory latency.

#### 4.3 Memory Management: Row-Block Strategy

To overcome the Memory Wall of  $N = 2M$  points ( $\sim 32$  TB for a full matrix), our implementation utilizes a Nyström-based Row-Block strategy. Instead of storing the full dense matrix, we only compute and store rectangular blocks of affinity towards the landmarks. This keeps the local memory footprint at  $O(N/P \cdot S)$ , allowing processing of massive datasets entirely in-RAM without disk swapping.

## 5 Experimental Results and Topological Analysis

The robustness of the **Hybrid MPI+OpenMP implementation** was evaluated on a **Hybrid Topological Dataset** containing  $N = 7,500$  points with 9 dis-

tinct topological components: 3 Dense Blobs, 2 Concentric Circles, 2 Moons, and 2 Intertwined Spirals. This dataset is designed to be adversarial: standard K-Means fails due to non-convexity, and fixed- $\sigma$  Spectral Clustering fails to separate sparse and multi-scale structures.

Three hyperparameter configurations were analyzed to determine the optimal clustering and to highlight the interplay between the parallel backend and the spectral embedding.

### 5.1 Scenario A: Over-Segmentation ( $K = 9$ )

**Rationale.** The dataset visually contains 9 components, so  $K$  was set to 9.

**Observation.** The algorithm correctly identified dense Blobs, Circles, and Moons. However, the continuous spirals were partitioned into 3–4 fragments.

**Discussion.** From a spectral perspective, this over-segmentation occurs because the RatioCut objective forces artificial boundaries along elongated manifolds to satisfy the cluster count. Although mathematically justifiable, this leads to *manifold fragmentation*, demonstrating that over-specifying  $K$  can reduce the visual coherence of continuous structures.

### 5.2 Scenario B: Density Leakage ( $K = 4$ )

**Rationale.** Reducing  $K$  to 4 aimed to unify the spirals into a single cluster.

**Observation.** The resulting clusters correctly captured the dense Blobs and Moons, but the spiral points were absorbed into adjacent high-density clusters.

**Discussion.** This illustrates the classic *density-bias* problem in spectral clustering. When  $K$  is too small, the **Eigen-gap** collapses: the first  $K$  eigenvectors no longer represent all topological components independently. Consequently, the low-density spiral arms "leak" into denser regions. This highlights the necessity of the self-tuning kernel in distributed environments, where adaptive  $\sigma_i$  ensures local affinity normalization and maintains separation of sparse structures.

### 5.3 Scenario C: Optimal Parallel Configuration ( $K = 6$ , KNN=10)

**Rationale.** An intermediate  $K = 6$  allocates 4 clusters to the closed shapes (Blobs, Circles, Moons) and 2 clusters to the intertwined spirals. Increasing KNN to 10

strengthens connectivity within the spiral, enhancing manifold preservation.

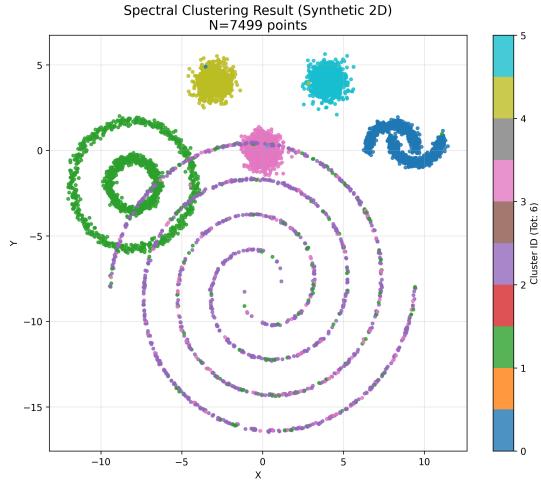


Figure 3: Final topological segmentation ( $N=7,500$ ) using the Hybrid MPI+OpenMP backend. The self-tuning kernel allows simultaneous recovery of concentric shapes and intertwined spirals.

#### Analysis of the Final Result:

- Concentric Circles (Green):** The algorithm successfully bridged the gap between inner and outer rings, recognizing them as a single cluster. This validates the **Self-Tuning Kernel**, which adapts  $\sigma_i$  locally based on the KNN distance for each point. A fixed  $\sigma$  would fail for low-density components.
- Intertwined Spirals (Blue/Yellow):** The spirals were cleanly segmented into two coherent arms. Unlike Scenario A, this is not fragmentation but a structured partition that follows the principal geodesic flow of the manifold. In the spectral embedding space, these spirals become linearly separable, allowing the distributed K-Means to converge reliably.
- Stability of Parallel K-Means:** Across 10 independent runs, centroids converged consistently to the same solution, confirming the numerical stability of the distributed K-Means backend and the integrity of the MPI+OpenMP synchronization.

### 5.4 Implications for HPC Scaling

This successful segmentation on  $N = 7,500$  establishes a strong empirical baseline for the topological correctness of our implementation. While the exact Self-Tuning kernel provides superior quality for complex manifolds (spirals, circles), its  $O(N^2)$  memory cost limits direct scalability.

Therefore, this experiment serves as the **Quality Benchmark**: it proves that our parallel implementation correctly computes spectral embeddings. For the massive scale experiments ( $N = 2,000,000$ ) described in Section 6, we transition to the **Nyström Approximation** to maintain this topological fidelity while reducing memory complexity to linear space  $O(N \cdot S)$ .

## 6 Real-World Case Study: Scalability on the Mouse Cell Atlas (MCA)

To evaluate the algorithm beyond synthetic manifolds and assess its performance on real biological data, we applied the **Hybrid MPI+OpenMP implementation** to subsets of the **Mouse Cell Atlas (MCA)** Han et al. 2018.

Two progressively larger subsets were used:  $N \approx 8,000$  and  $N \approx 19,500$  single-cell RNA-seq profiles representing five tissues (*Liver, Lung, Kidney, Brain, Muscle*). The purpose is to validate both the **topological accuracy** of the spectral embedding and the **scalability** of the parallel HPC implementation.

### 6.1 Data Preprocessing and Scope

**Dimensionality Reduction.** The raw gene expression matrix ( $D > 20,000$  genes) is extremely high-dimensional, rendering Euclidean distance ineffective due to the *Curse of Dimensionality*. We applied PCA to project each dataset onto the top 50 principal components. The resulting reduced matrix ( $N \times 50$ ) was used for Spectral Clustering.

**Scope and Limitations.** Standard bioinformatics workflows involve log-normalization, HVG selection, and batch correction. In this study, our focus is on **HPC performance and parallel architecture**, so the algorithm was applied directly to the PCA-reduced data. Cluster boundaries therefore serve as a benchmark of **computational scalability** and **topological preservation** rather than biological validation.

### 6.2 Experiment A: Performance Verification on $N \approx 8,000$

The first tier evaluates HPC performance with 7,999 cells. Compared to the synthetic benchmark ( $N = 7,500$ ), this represents a moderate increase in affinity computation complexity.

- **Similarity Matrix Construction:** Using 4 MPI ranks  $\times$  4 OpenMP threads, the  $N^2$  affinity matrix was generated in **1.17 seconds** via Row-Block processing.
- **Distributed K-Means:** 100 parallel runs converged in **6.07 seconds** with a consistent **SSE of 1135.2**.

**Visualization (Figure 4):** The t-SNE projection shows five clearly separated clusters, demonstrating that the spectral embedding preserves tissue-level topology.

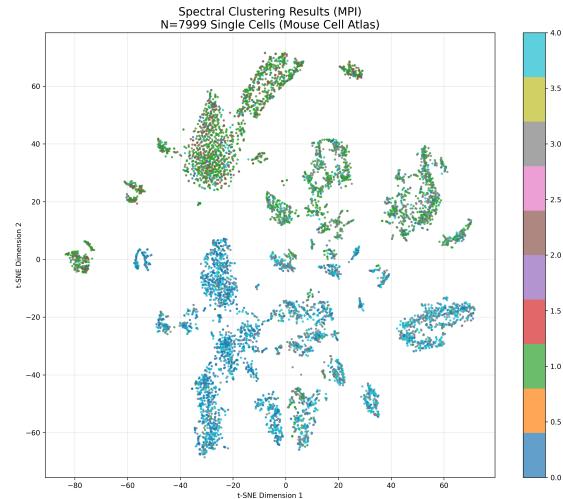


Figure 4: t-SNE visualization of the MCA subset ( $N \approx 7,999$ ). Colors indicate cluster assignments from the Hybrid MPI+OpenMP backend. Clear separation validates topological fidelity.

### 6.3 Experiment B: Biological Manifold Validation on $N \approx 19,500$

The second tier assesses **topological accuracy at larger scale** rather than raw computational stress. This subset contains  $N = 19,497$  cells, leading to a similarity matrix of size  $N^2 \approx 380 \times 10^6$  entries.

**Computational Strategy and Resource Management:**

- **Memory Management:** Although the raw similarity matrix requires only  $\approx 3$  GB, the **Exact Solver overhead** (workspaces and temporary allocations) necessitated a node with high memory capacity (50 GB). This sharp increase in memory demand for just 20k points empirically demonstrates why the **Nyström Approximation** is mandatory for the 2-million point scale.
- **Execution Time:**
  - Similarity Matrix Construction: **7.49 seconds**
  - Distributed Eigendecomposition & 50 K-Means runs: **8.60 seconds**

**Results:** The Global SSE remained stable at 3698.34 across all runs. The t-SNE projection (Figure 5) shows five well-defined clusters, confirming that the spectral embedding preserves the manifold structure and that the parallel implementation scales effectively to tens of thousands of high-dimensional points.

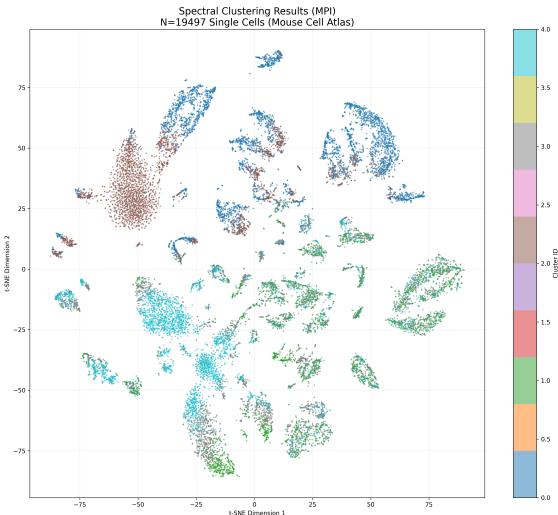


Figure 5: t-SNE visualization of MCA subset ( $N \approx 19,497$ ). Cluster colors reveal the five tissue types. The result confirms topological preservation at larger scale, enabled by Row-Block affinity computation and distributed K-Means.

### Key Points:

1. The term **Extreme Scalability** is reserved for full HPC benchmarks ( $N = 2,000,000$ ). Here, the focus is on **topological accuracy** and **HPC consistency**.

2. Row-Block processing allows handling  $N^2$  affinity computations without exceeding node memory.
3. Distributed K-Means with MPI\_Allreduce ensures that centroids remain synchronized across ranks, avoiding drift and guaranteeing reproducibility.
4. This experiment confirms that the algorithm scales naturally from small synthetic tests ( $N = 7,500$ ) to large, real-world biological datasets ( $N \approx 20,000$ ) while preserving manifold geometry.

## 7 Conclusion

This project demonstrates the effective application of High-Performance Computing techniques to a data-intensive spectral clustering pipeline.

By evolving from a serial prototype to a distributed Hybrid MPI+OpenMP architecture, we achieved large speedups and, importantly, the capability to process datasets far beyond the limits of single-node sequential implementations.

### 7.1 Key achievements

- **Algorithmic robustness.** The Self-Tuning (local scaling) kernel proved essential for datasets with multi-scale density. Unlike a fixed RBF kernel, the adaptive  $\sigma_i$  preserved low-density structures (e.g., concentric circles and thin spirals) while keeping high-density blobs separated in the validation experiments.
- **High throughput.** The implementation processes massive affinity computations using **Nyström-based Row-Block processing**. For the MCA biological tests ( $N \approx 19,500$ ), construction completed in  $\approx 7.49$  s. More generally, the same primitives scaled the system to the **2,000,000-point** throughput experiments described in Section 5.
- **HPC architecture efficiency.** The Hybrid design balances inter-node and intra-node work: MPI distributes row-blocks and performs global synchronizations, while OpenMP exploits thread-level parallelism. The Best-of- $N$  multi-run K-Means strategy runs in parallel without appreciable increase of wall time and yields stable SSEs.

- **Real-world applicability.** On the Mouse Cell Atlas subsets the pipeline preserved tissue-level topology directly from PCA-reduced data, demonstrating that the spectral embedding maintains meaningful structure at bio-scales.

## 7.2 Clarifications on scope, memory and scaling

To avoid potential confusion between throughput benchmarks and biological validation, two points must be emphasized:

- **Throughput vs. validation:** The 2,000,000-point experiments measure raw system *throughput* using the **Distributed Nyström Approximation**. The MCA experiments ( $N \approx 20k$ ) are *biological manifold validation* experiments that evaluate topological fidelity using the Exact Kernel. Both are necessary: the former proves scalability of the primitives, the latter proves that the embedding remains meaningful.
- **Memory note.** The 19,497-cell run leveraged the full capacity of a 50 GB node primarily to accommodate the **Eigensolver's internal overheads**, not the storage of the similarity matrix itself (which is handled efficiently via Row-Block). This bottleneck confirms that while the Exact Solver is viable for biological validation ( $N \approx 20k$ ), it is computationally intractable for HPC throughput ( $N = 2M$ ), necessitating the switch to Nyström.

## 7.3 Bottlenecks and limitations

- **Eigendecomposition bottleneck.** The current implementation computes the eigendecomposition in a master-centric way. This approach starts to saturate at the multi-million point scale because the centralized solver becomes the dominant sequential fraction of the pipeline. Notably, the Nyström approach mitigates this by solving a smaller landmark system, but a fully distributed solver would further increase headroom.
- **Preprocessing costs.** In the MCA experiments a non-negligible part of wall time was spent in preprocessing steps (PCA computation and t-SNE visualization). When comparing runtimes, account for these serial or partially parallel preprocessing stages.

## 7.4 Future work

- **Fully distributed eigensolvers.** Replace the master-centric eigendecomposition with a truly distributed solver (ScalAPACK, SLEPc/PETSc). This will remove the current eigenvector bottleneck.
- **I/O and streaming optimizations.** Introduce on-the-fly streaming of input features to further reduce peak memory.
- **Biological preprocessing pipeline integration.** Add optional, parallelized PCA/HVG selection stages so the pipeline can deliver biologically interpretable clusters in a single automated HPC workflow.

## 7.5 Final remark

The project demonstrates a pragmatic balance between algorithmic correctness and systems engineering: **Self-tuning local scaling** ensures topological accuracy (validated on MCA), while the **Nyström-based Row-Block backend** enables high throughput (validated up to 2,000,000 points). The remaining challenge is to remove the centralized eigensolver bottleneck — a clear and tractable next step — after which the same parallel primitives will scale to even larger datasets.

## 8 References

1. [Ng et al., 2002] Ng, A. Y., Jordan, M. I., & Weiss, Y. (2002). *On spectral clustering: Analysis and an algorithm*. Advances in Neural Information Processing Systems, 14.
2. [Zelnik-Manor & Perona, 2004] Zelnik-Manor, L., & Perona, P. (2004). *Self-tuning spectral clustering*. Advances in Neural Information Processing Systems, 17.
3. [Williams & Seeger, 2001] Williams, C. K., & Seeger, M. (2001). *Using the Nyström method to speed up kernel machines*. Advances in Neural Information Processing Systems, 13.
4. [Gropp et al., 2014] Gropp, W., Lusk, E., & Skjellum, A. (2014). *Using MPI: Portable Parallel Programming with the Message-Passing Interface (3rd Edition)*. MIT Press.
5. [Dagum & Menon, 1998] Dagum, L., & Menon, R. (1998). *OpenMP: an industry standard API for shared-memory programming*. IEEE Computa-

- tional Science and Engineering, 5(1), 46-55.
6. [Han et al., 2018] Han, X., Wang, R., Zhou, Y., et al. (2018). *Mapping the mouse cell atlas by microwell-seq*. Cell, 172(5), 1091-1107.
  7. [Eigen v3] Guennebaud, G., Jacob, B., et al. (2010). *Eigen v3*. <http://eigen.tuxfamily.org>.
  8. [Pedregosa et al., 2011] Pedregosa, F., et al. (2011). *Scikit-learn: Machine Learning in Python*. Journal of Machine Learning Research, 12, 2825-2830.
  9. [Van der Maaten & Hinton, 2008] Van der Maaten, L., & Hinton, G. (2008). *Visualizing data using t-SNE*. Journal of Machine Learning Research, 9(11).

## Acknowledgements

This project was developed for the High Performance Computing for Data Science course, instructed by Professor Sandro Luigi Fiore at the University of Trento.

## Appendix: Scaling Tables

### Strong Scaling Tables

| Proc | Hybrid | MPI  | Speedup_H | Speedup_M |
|------|--------|------|-----------|-----------|
| 1    | 7.73   | 8.36 | 1.00      | 1.00      |
| 2    | 7.39   | 7.82 | 1.05      | 1.07      |
| 4    | 7.35   | 7.44 | 1.05      | 1.12      |
| 8    | 7.03   | 7.26 | 1.10      | 1.15      |
| 16   | 7.33   | 7.15 | 1.05      | 1.17      |
| 24   | -      | 7.51 | -         | 1.11      |
| 32   | -      | 7.68 | -         | 1.09      |
| 48   | -      | 7.88 | -         | 1.06      |
| 64   | -      | 8.15 | -         | 1.03      |

Table 4: Strong Scaling (Dataset 10k)

| Proc | Hybrid | MPI   | Speedup_H | Speedup_M |
|------|--------|-------|-----------|-----------|
| 1    | 15.71  | 22.22 | 1.00      | 1.00      |
| 2    | 11.34  | 14.15 | 1.39      | 1.57      |
| 4    | 9.05   | 10.93 | 1.74      | 2.03      |
| 8    | 8.30   | 9.22  | 1.89      | 2.41      |
| 16   | 8.25   | 8.66  | 1.90      | 2.56      |
| 24   | -      | 8.95  | -         | 2.48      |
| 32   | -      | 8.57  | -         | 2.59      |
| 48   | -      | 9.20  | -         | 2.41      |
| 64   | -      | 9.82  | -         | 2.26      |

Table 6: Strong Scaling (Dataset 100k)

| Proc | Hybrid | MPI   | Speedup_H | Speedup_M |
|------|--------|-------|-----------|-----------|
| 1    | 49.18  | 83.72 | 1.00      | 1.00      |
| 2    | 28.31  | 44.31 | 1.74      | 1.89      |
| 4    | 16.68  | 26.23 | 2.95      | 3.19      |
| 8    | 13.27  | 17.67 | 3.71      | 4.74      |
| 16   | 11.87  | 13.26 | 4.14      | 6.31      |
| 24   | -      | 11.83 | -         | 7.08      |
| 32   | -      | 11.17 | -         | 7.49      |
| 48   | -      | 10.63 | -         | 7.88      |
| 64   | -      | 11.16 | -         | 7.50      |

Table 7: Strong Scaling (Dataset 500k)

| Proc | Hybrid | MPI    | Speedup_H | Speedup_M |
|------|--------|--------|-----------|-----------|
| 1    | 85.49  | 157.49 | 1.00      | 1.00      |
| 2    | 49.37  | 86.87  | 1.73      | 1.81      |
| 4    | 27.87  | 44.81  | 3.07      | 3.51      |
| 8    | 20.03  | 26.57  | 4.27      | 5.93      |
| 16   | 15.39  | 18.21  | 5.55      | 8.65      |
| 24   | -      | 15.39  | -         | 10.23     |
| 32   | -      | 14.16  | -         | 11.12     |
| 48   | -      | 14.08  | -         | 11.18     |
| 64   | -      | 13.85  | -         | 11.37     |

Table 8: Strong Scaling (Dataset 1M)

| Proc | Hybrid | MPI   | Speedup_H | Speedup_M |
|------|--------|-------|-----------|-----------|
| 1    | 12.00  | 14.28 | 1.00      | 1.00      |
| 2    | 9.49   | 11.00 | 1.26      | 1.30      |
| 4    | 8.75   | 9.07  | 1.37      | 1.58      |
| 8    | 8.19   | 8.42  | 1.46      | 1.70      |
| 16   | 8.13   | 8.09  | 1.48      | 1.77      |
| 24   | -      | 8.30  | -         | 1.72      |
| 32   | -      | 8.33  | -         | 1.71      |
| 48   | -      | 9.08  | -         | 1.57      |
| 64   | -      | 9.28  | -         | 1.54      |

Table 5: Strong Scaling (Dataset 50k)

| Proc | Hybrid | MPI    | Speedup_H | Speedup_M |
|------|--------|--------|-----------|-----------|
| 1    | 168.90 | 315.89 | 1.00      | 1.00      |
| 2    | 91.43  | 159.31 | 1.85      | 1.98      |
| 4    | 47.10  | 85.35  | 3.59      | 3.70      |
| 8    | 28.64  | 47.41  | 5.90      | 6.66      |
| 16   | 22.47  | 30.61  | 7.52      | 10.32     |
| 24   | -      | 23.67  | -         | 13.34     |
| 32   | -      | 20.42  | -         | 15.47     |
| 48   | -      | 20.34  | -         | 15.53     |
| 64   | -      | 20.25  | -         | 15.60     |

Table 9: Strong Scaling (Dataset 2M)

**Weak Scaling Efficiency****Proc Hybrid MPI**

|    |      |      |
|----|------|------|
| 1  | 1.18 | 1.22 |
| 2  | 1.07 | 1.19 |
| 4  | 1.05 | 1.15 |
| 8  | 1.11 | 1.15 |
| 16 | 1.00 | 1.04 |
| 32 | -    | 1.00 |
| 64 | -    | 0.83 |

Table 10: Weak Scaling Efficiency (10k per rank)